

**Московский государственный технический
университет им. Н. Э. Баумана**

**Курс «Технологии машинного обучения»
Отчёт по лабораторной работе №5
«Ансамбли моделей машинного обучения»**

Выполнила:
Шимолина П.К.,
группа ИУ5-61Б

Проверил:
Нардид А.Н.,
каф. ИУ5

Дата:

Дата:

Подпись:

Подпись:

2023
Москва

lab-5-2

June 14, 2023

```
[1]: import pandas as pd
import numpy as np
import plotly.express as px
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn
import category_encoders as ce
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from mlxtend.regressor import StackingCVRegressor
from sklearn.datasets import make_regression
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LinearRegression
```

```
[2]: df = pd.read_csv("Life Expectancy Data.csv")
```

```
[3]: y = df["Life expectancy "]
X = df.drop(["Life expectancy "], axis=1)
```

```
[4]: y.fillna(y.median(), inplace=True)
```

```
[5]: X.Year = pd.to_datetime(X.Year).dt.year
```

```
[6]: bin_enc = ce.BinaryEncoder(drop_invariant=True)
X = bin_enc.fit_transform(X)
```

```
[7]: X.fillna(X.mean(), inplace=True)
```

```
[8]: X_train, X_test, y_train, y_test= train_test_split(X, y, test_size= 0.30,
↳random_state=9)
```

0.1 ()

```
[9]: rfr = RandomForestRegressor()
rfr.fit(X_train, y_train)
prediction = rfr.predict(X_test)
```

```
[10]: params = {
    'max_depth' : [10, 15, 20],
}

grid = GridSearchCV(estimator=RandomForestRegressor(), param_grid=params, cv= 5)
grid.fit(X_train, y_train)
```

```
[10]: GridSearchCV(cv=5, estimator=RandomForestRegressor(),
    param_grid={'max_depth': [10, 15, 20]})
```

```
[11]: grid.best_score_, grid.best_params_
```

```
[11]: (0.9560259403650033, {'max_depth': 20})
```

```
[12]: rfr = RandomForestRegressor(max_depth=15, random_state=0)
rfr.fit(X_train, y_train)
prediction = rfr.predict(X_test)
mean_squared_error(y_test, prediction)
```

```
[12]: 3.780427824778443
```

0.2

```
[13]: params = {
    'n_estimators': [500, 800],
    'max_depth': [5, 8],
    'min_samples_split': [2, 5],
    'learning_rate': [0.01, 0.1]
}
grid = GridSearchCV(GradientBoostingRegressor(), param_grid=params, cv=5,
    ↪n_jobs=-1)
grid.fit(X_train, y_train)
```

```
[13]: GridSearchCV(cv=5, estimator=GradientBoostingRegressor(), n_jobs=-1,
    param_grid={'learning_rate': [0.01, 0.1], 'max_depth': [5, 8],
    'min_samples_split': [2, 5],
    'n_estimators': [500, 800]})
```

```
[14]: grid.best_score_, grid.best_params_
```

```
[14]: (0.9572816094106351,  
      {'learning_rate': 0.1,  
       'max_depth': 5,  
       'min_samples_split': 5,  
       'n_estimators': 800})
```

```
[15]: gbr = GradientBoostingRegressor(**grid.best_params_)  
      gbr.fit(X_train, y_train)  
      prediction = gbr.predict(X_test)
```

```
[16]: mean_squared_error(y_test, prediction)
```

```
[16]: 3.428038590524013
```

0.3

```
[17]: reg1 = RandomForestRegressor(random_state=42)  
      reg2 = GradientBoostingRegressor(random_state=42)  
      reg3 = LinearRegression()  
      meta_learner = LinearRegression()  
      sr = StackingCVRegressor(regressors=[reg1, reg2, reg3],  
                               ↪meta_regressor=meta_learner)  
      params = {  
          'randomforestregressor__max_depth': [5, 8]  
      }  
  
      grid = GridSearchCV(estimator=sr, param_grid=params, cv=5, n_jobs=-1)  
      grid.fit(X_train, y_train)
```

```
[17]: GridSearchCV(cv=5,  
                 estimator=StackingCVRegressor(meta_regressor=LinearRegression(),  
                 regressors=[RandomForestRegressor(random_state=42),  
                 GradientBoostingRegressor(random_state=42),  
                 LinearRegression()])),  
                 n_jobs=-1,  
                 param_grid={'randomforestregressor__max_depth': [5, 8]})
```

```
[18]: sr.get_params().keys()
```

```
[18]: dict_keys(['cv', 'meta_regressor__copy_X', 'meta_regressor__fit_intercept',  
             'meta_regressor__n_jobs', 'meta_regressor__normalize',
```

```

'meta_regressor__positive', 'meta_regressor', 'multi_output', 'n_jobs',
'pre_dispatch', 'random_state', 'refit', 'regressors', 'shuffle',
'store_train_meta_features', 'use_features_in_secondary', 'verbose',
'randomforestregressor', 'gradientboostingregressor', 'linearregression',
'randomforestregressor__bootstrap', 'randomforestregressor__ccp_alpha',
'randomforestregressor__criterion', 'randomforestregressor__max_depth',
'randomforestregressor__max_features', 'randomforestregressor__max_leaf_nodes',
'randomforestregressor__max_samples',
'randomforestregressor__min_impurity_decrease',
'randomforestregressor__min_samples_leaf',
'randomforestregressor__min_samples_split',
'randomforestregressor__min_weight_fraction_leaf',
'randomforestregressor__n_estimators', 'randomforestregressor__n_jobs',
'randomforestregressor__oob_score', 'randomforestregressor__random_state',
'randomforestregressor__verbose', 'randomforestregressor__warm_start',
'gradientboostingregressor__alpha', 'gradientboostingregressor__ccp_alpha',
'gradientboostingregressor__criterion', 'gradientboostingregressor__init',
'gradientboostingregressor__learning_rate', 'gradientboostingregressor__loss',
'gradientboostingregressor__max_depth',
'gradientboostingregressor__max_features',
'gradientboostingregressor__max_leaf_nodes',
'gradientboostingregressor__min_impurity_decrease',
'gradientboostingregressor__min_samples_leaf',
'gradientboostingregressor__min_samples_split',
'gradientboostingregressor__min_weight_fraction_leaf',
'gradientboostingregressor__n_estimators',
'gradientboostingregressor__n_iter_no_change',
'gradientboostingregressor__random_state',
'gradientboostingregressor__subsample', 'gradientboostingregressor__tol',
'gradientboostingregressor__validation_fraction',
'gradientboostingregressor__verbose', 'gradientboostingregressor__warm_start',
'linearregression__copy_X', 'linearregression__fit_intercept',
'linearregression__n_jobs', 'linearregression__normalize',
'linearregression__positive'])

```

```
[19]: grid.best_score_, grid.best_params_
```

```
[19]: (0.9501129846734061, {'randomforestregressor__max_depth': 8})
```

```
[20]: reg1 = RandomForestRegressor(random_state=42, max_depth=8)
```

```
[21]: sr = StackingCVRegressor(regressors=[reg1, reg2, reg3],
    ↪ meta_regressor=meta_learner)
sr.fit(X_train, y_train)
prediction = sr.predict(X_test)
```

```
C:\anaconda3\lib\site-packages\sklearn\base.py:443: UserWarning: X has feature names, but RandomForestRegressor was fitted without feature names
```

```
warnings.warn(
```

```
C:\anaconda3\lib\site-packages\sklearn\base.py:443: UserWarning: X has feature names, but GradientBoostingRegressor was fitted without feature names
```

```
warnings.warn(
```

```
C:\anaconda3\lib\site-packages\sklearn\base.py:443: UserWarning: X has feature names, but LinearRegression was fitted without feature names
```

```
warnings.warn(
```

```
[22]: mean_squared_error(y_test, prediction)
```

```
[22]: 4.43231527722557
```